

Last time

- Fully connected layer
- Activation functions

Today

- convolutional (2D) layer
- pooling layers

Ex. $P = 9, d_1 = 16, d_2 = 8, d_3 = 1$

$$f_{\text{dense}}(x) = [\max(0, \max(0, xW_1 + B_1)W_2 + B_2)]W_3 + B_3$$

$$x = [n, P]$$

$$\text{Total parameter count} = 305$$

$$x_i = [1, P]$$

$$\text{1st layer parameter count} = 160$$

Now

$$x = [n, \sqrt{P}, \sqrt{P}, 1]$$

$$x_i = \begin{bmatrix} 1 & & & & 3 \\ \hline & \diagup & \diagup & \diagup & \\ & \diagdown & \diagdown & \diagdown & \\ & \diagup & \diagup & \diagup & \\ & \diagdown & \diagdown & \diagdown & \\ \hline & & & & 3 \end{bmatrix}$$

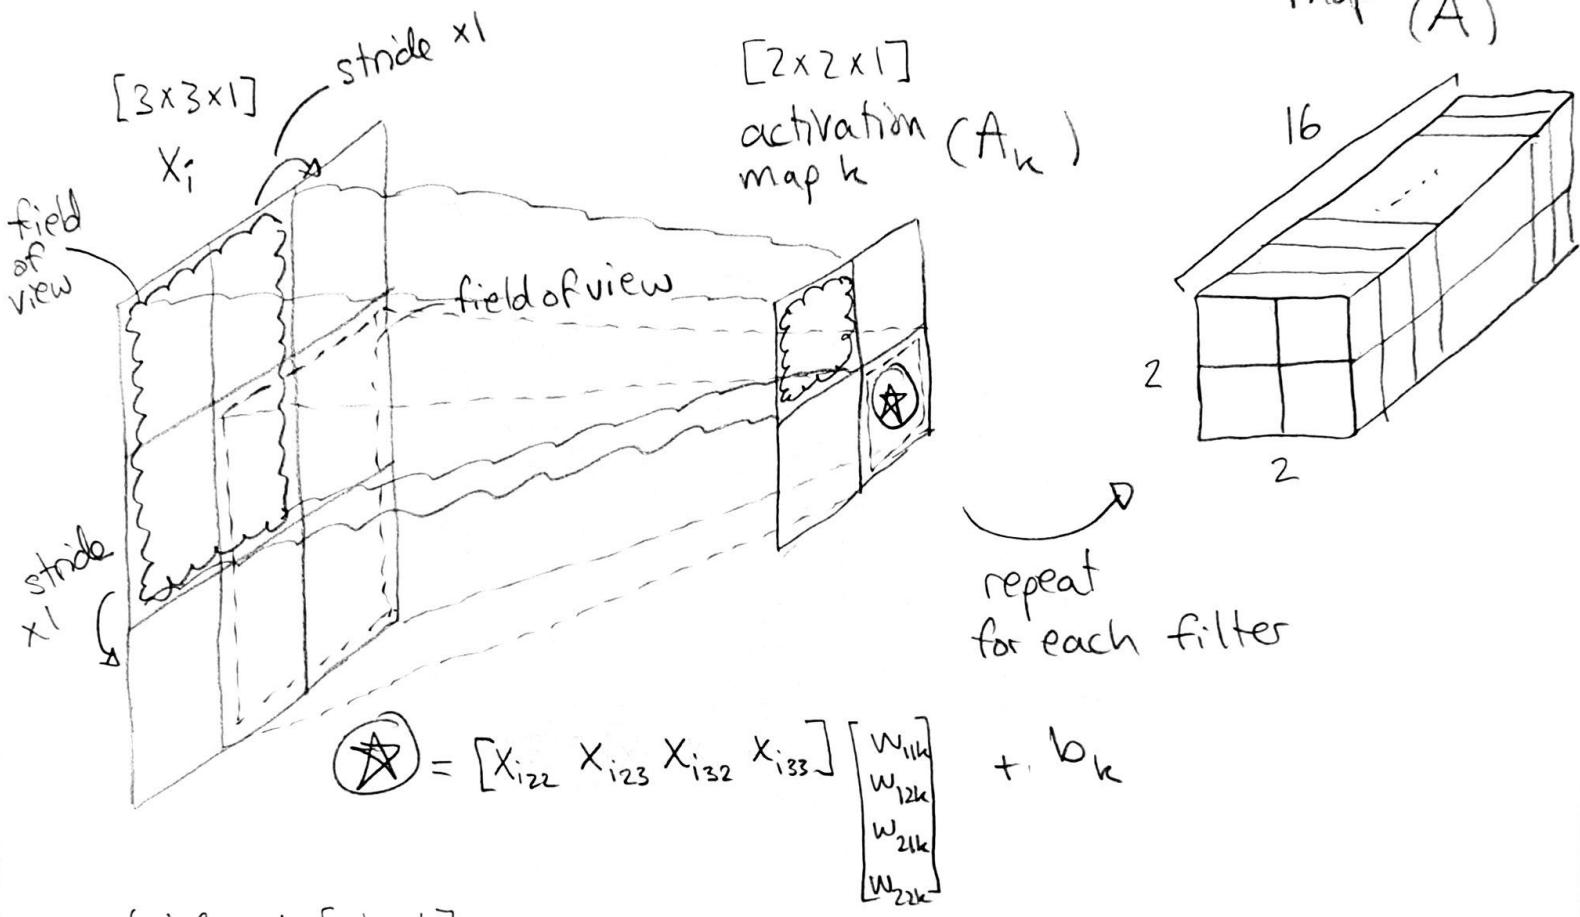
$$x_i = [1, \sqrt{P}, \sqrt{P}, 1]$$

$$f_{\text{conv}}(x) = [\max(0, \text{flatten}(\text{Conv2D}(x))W_2 + B_2)]W_3 + B_3$$

2D Convolutional Layer

- a restricted type of template checking
- assumes spatial structure, templates are only checked against a spatially-restricted window referred to as the field of view.
- commonly used to process image data
- filter: ^(or kernel) the template being checked for. Analogous to a neuron in the dense layer.
 - W : [filter height, filter width, filter depth]
 - B : [1]
- Note: layer's parameter size is not a function of p (the input size)
- strides: the step sizes, in both the height and width directions, that are used to move the field of view across the input.
- activation map: the spatially organized output of the application of a filter to the input.

Note: Conv layer is not fully connected because each element of an input is not considered in each "check" of the template.



strides : $[1 \times 1]$

filter size : $[2 \times 2 \times 1]$ ($\begin{bmatrix} w_{11k} & w_{12k} \\ w_{21k} & w_{22k} \end{bmatrix}, b_k$)

filters : 16

The output is an "image" of size $[2 \times 2 \times 16]$

$$\text{Conv2D}(X) = \sigma(A(X))$$

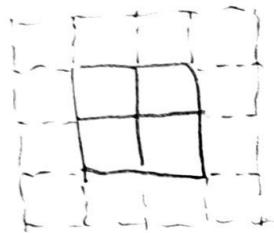
Notes:

- The number of parameters in this layer ($80 = (4+1) \times 16$) is much smaller than the Dense layer used in the example last time which contained the same number of templates (16) and operated on the same sized input (i.e. 1×9 vs $1 \times \sqrt{9} \times \sqrt{9} \times 1$).
- The total parameter size of $f_{\text{conv}}(x)$:
$$\begin{aligned} &= 80 + (16 \times 4 \times 8) + 8 + (8 \times 1) + 1 \\ &= 80 + 520 + 9 = 609 \Leftrightarrow f_{\text{dense}} \end{aligned}$$
- The parameter count for a 2D Conv layer:
$$= (\text{filter height} \times \text{filter width} \times \text{filter depth} + 1) \times \# \text{filters}$$
- Filter size and strides has a big influence on output size.
→ output size = $1 + \left\lceil \frac{\text{input size} - \text{filter size} + 2 \times \text{padding}}{\text{stride}} \right\rceil$
↑ for each dimension (ie height and width)

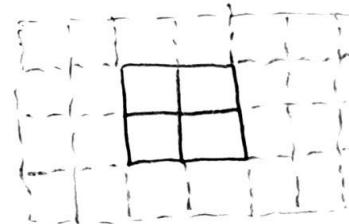
Padding: the concatenation of additional height or width elements to the input before template checks

e.g.

$$\text{padding} = [1 \times 1]$$



$$\text{padding} = [2 \times 1]$$

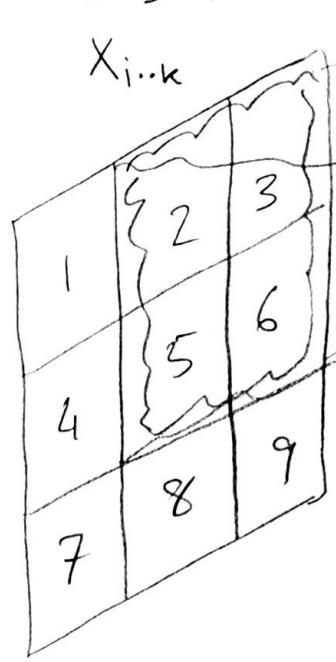


- used to evenly apply filters to each element of the input
- padded elements are usually zero or the mean of mirror-flipping of adjacent elements.

Pooling Layers (2D)

- used to downsample / summarize inputs
- (usually) contain no learnable parameters.
- characterized by their pooling operation (g)
 - max, mean
- operate similarly to conv2D layers and are also thus parameterized by filter size, padding, and strides.

$3 \times 3 \times 1$



A_k

$$g\left(x_{i12k}, x_{i13k}, x_{i22k}, x_{i23k}\right)$$

if $g = \max$

then

$$A_k = \begin{array}{|c|c|} \hline 5 & 6 \\ \hline 8 & 9 \\ \hline \end{array}$$

if $g = \text{mean}$

strides : $[1 \times 1]$

filter size : $[2 \times 2]$

$$A_k = \begin{array}{|c|c|} \hline 3 & 4 \\ \hline 6 & 7 \\ \hline \end{array}$$

→ Note : Unlike convolutional layers, Pooling layers operate on a single channel at a time.