

Last time

- convolutional layers
 - pooling layers
 - smaller parameter counts vs dense layers
- } capturing spatial structures

Today

- other important types of layers
 - [→ Batch Normalization
 - Dropout
 - [→ Residual Layer
 - can learn to be removed during training
- } different behavior at train & test time

Batch Normalization

- a "special" layer type which does not follow the linear + nonlinear piece recipe
- is usually applied either: (1) after the linear piece and before the non-linear piece of another layer, or (2) after the non-linear piece of another layer.
 - (1) $\sigma(\text{BatchNorm}(z(x)))$
 - (2) $\text{BatchNorm}(\sigma(z(x)))$
- was introduced as a fix to the gradient explosion / gradient vanishing problem, but is perhaps best thought of as a noising / regularizing layer.
- Idea: center and scale the activations of a layer to prevent the drift of this distribution to extremes during training OR reparameterize the activations of each observation as measuring the relative differences between A and members of the batch.

input: $X = [n, p]$ e.g. \leftarrow output of previous layer
 batch size

Behavior during training:

(1) Obtain the batch mean \bar{X}_B and batch variance s_B^2 feature-wise:

$$\bar{X}_B = \frac{1}{n} \sum_{i=1}^n x_i \in \mathbb{R}^p$$

⚠ highly dependent
on batch size

$$s_B^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{X}_B)^2$$

(2) Update the moving average estimators of the population mean and variance of activations:

$$\hat{\mu} = \lambda \hat{\mu} + (1-\lambda) \bar{X}_B$$

$$\hat{\sigma}^2 = \lambda \hat{\sigma}^2 + (1-\lambda) s_B^2$$

λ is the momentum

hyperparameter (~ 0.99)

(3) calculate normalized activations:

$$\tilde{x}_i = \left(\frac{x_i - \bar{X}_B}{\sqrt{s_B^2 + \epsilon}} \right)$$

(4) learned scale and shift :

$$z_i = \gamma \tilde{x}_i + \beta$$

where $\gamma \in \mathbb{R}^P$, $\beta \in \mathbb{R}^P$ are trainable parameters, initialized to 1 and 0 respectively

→ allows the network to "undo" the scaling

Behavior during testing :

$$z_i = \gamma \left(\frac{x_i - \hat{\mu}}{\sqrt{\hat{\sigma}^2 + \epsilon}} \right) + \beta$$

→ replace batch mean and variance estimates with estimates of population statistics.

→ why moving average? distribution may change over time during training

Related layers

→ Layer Norm

Dropout

- a "special" layer that does not have the usual linear → nonlinear structure
- is usually applied before a Batch Normalization layer e.g. $\text{BatchNorm}(\text{Dropout}(\sigma(z(x))))$
- Idea: Noise a layer's outputs by "dropping" each neuron (i.e. setting its value to zero) for each observation with a probability d .

Behavior during training:

$$\text{Dropout}_d(X) = X \cdot M$$

$[n \times p] \quad [n \times p]$

element-wise multiplication

$$M_{ij} = \mathbb{1}\{U_{ij} > d\}, \quad U_{ij} \sim \text{uniform}(0, 1)$$

Behavior during testing:

$$\text{Dropout}_d(X) = X \cdot (1-d)$$

- want the expected magnitudes of the outputs to be the same during training and testing

Behavior during training (v2):

$$\text{Dropout}_d(X) = X \cdot M / (1 - d)$$

Behavior during testing (v2):

$$\text{Dropout}_d(X) = X$$

→ avoids one computation during inference (where latency can be a big issue)


Related: Stochastic Depth (layer)
Drop Connect (weight drop)

→ Interpretations:

→ prevent over-reliance on any particular feature


→ encourage utilization of all features

→ learn features that are useful for all sub-networks of the trained model



Residual Layer

→ Idea: Enable a model to choose/learn the optimal number of layers needed during training.



$$Z_2 = Z_1 + X \quad (\text{where } + \text{ is element-wise addition})$$

→ Note: Z_1 and X must have the same shape

→ A key innovation which enabled training of very deep (> 100 layer) networks, overcoming vanishing/exploding gradient problem.